

Weather Prediction

Will it rain Tomorrow?

Muhammed Berk Önder

Istanbul Data Science Academy

January 21, 2022

Will it rain tomorrow?

- **Problem:** Can we tell from today that it will rain tomorrow?
- **Whom does it concern:** Anyone who doesn't want to get wet tomorrow.

- Pull sql table from database with *sqlalchemy*.
- Analyse and Clean the data.
- Used *RandomForestClassifier* as a model.
- Deploy with *Flask*.

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 145460 entries, 0 to 145459
```

```
Data columns (total 23 columns):
```

#	Column	Non-Null	Count	Dtype
0	date	145460	non-null	datetime64[ns]
1	location	145460	non-null	object
2	mintemp	143975	non-null	float64
3	maxtemp	144199	non-null	float64
4	rainfall	142199	non-null	float64
5	evaporation	82670	non-null	float64
6	sunshine	75625	non-null	float64
7	windgustdir	135134	non-null	object
8	windgustspeed	135197	non-null	float64
9	winddir9am	134894	non-null	object
10	winddir3pm	141232	non-null	object
11	windspeed9am	143693	non-null	float64
12	windspeed3pm	142398	non-null	float64
13	humidity9am	142806	non-null	float64
14	humidity3pm	140953	non-null	float64
15	pressure9am	130395	non-null	float64
16	pressure3pm	130432	non-null	float64
17	cloud9am	89572	non-null	float64
18	cloud3pm	86102	non-null	float64
19	temp9am	143693	non-null	float64
20	temp3pm	141851	non-null	float64
21	raintoday	142199	non-null	object
22	raintomorrow	142193	non-null	object

```
dtypes: datetime64[ns](1), float64(16), object(6)
```


```
memory usage: 25.5+ MB
```





```
df.dtypes
```





```
date                datetime64[ns]
location            object
mintemp             float64
maxtemp             float64
rainfall            float64
evaporation         float64
sunshine            float64
windgustdir         object
windgustspeed       float64
winddir9am          object
winddir3pm          object
windspeed9am        float64
windspeed3pm        float64
humidity9am          float64
humidity3pm          float64
pressure9am          float64
pressure3pm          float64
cloud9am            float64
cloud3pm            float64
temp9am             float64
temp3pm             float64
raintoday           object
raintomorrow        object
dtype: object
```


```
df.isnull().sum()
```

```
date            0
location        0
mintemp         1485
maxtemp         1261
rainfall        3261
evaporation     62790
sunshine        69835
windgustdir     10326
windgustspeed   10263
winddir9am      10566
winddir3pm      4228
windspeed9am    1767
windspeed3pm    3062
humidity9am     2654
humidity3pm     4507
pressure9am     15065
pressure3pm     15028
cloud9am        55888
cloud3pm        59358
temp9am         1767
temp3pm         3609
raintoday       3261
raintomorrow    3267
dtype: int64
```










Correlation Heatmap of Rain in Australia Dataset



From the above correlation heat map, we can conclude that:

- MinTemp and MaxTemp variables are highly positively correlated (correlation coefficient = 0.74).
- MinTemp and Temp3pm variables are also highly positively correlated (correlation coefficient = 0.71).
- MinTemp and Temp9am variables are strongly positively correlated (correlation coefficient = 0.90).
- MaxTemp and Temp9am variables are strongly positively correlated (correlation coefficient = 0.89).
- MaxTemp and Temp3pm variables are also strongly positively correlated (correlation coefficient = 0.98).
- WindGustSpeed and WindSpeed3pm variables are highly positively correlated (correlation coefficient = 0.69).
- Pressure9am and Pressure3pm variables are strongly positively correlated (correlation coefficient = 0.96).
- Temp9am and Temp3pm variables are strongly positively correlated (correlation coefficient = 0.86).

```
def max_value(df4, variable, top):
    return np.where(df4[variable]>top, top, df4[variable])

for df4 in [X_train, X_test]:
    df4['rainfall'] = max_value(df4, 'rainfall', 3.2)
    df4['evaporation'] = max_value(df4, 'evaporation', 21.8)
    df4['windspeed9am'] = max_value(df4, 'windspeed9am', 55)
    df4['windspeed3pm'] = max_value(df4, 'windspeed3pm', 57)
```

```
print('Rainfall:', X_train.rainfall.max(), X_test.rainfall.max())
print('Evaporation:', X_train.evaporation.max(), X_test.evaporation.max())
print('WindSpeed9am:', X_train.windspeed9am.max(), X_test.windspeed9am.max())
print('WindSpeed3pm:', X_train.windspeed3pm.max(), X_test.windspeed3pm.max())

print(X_train[numerical].describe())
```

Rainfall: 3.2 3.2

Evaporation: 21.8 21.8

WindSpeed9am: 55.0 55.0

WindSpeed3pm: 57.0 57.0

	mintemp	maxtemp	rainfall	evaporation \
count	116368.000000	116368.000000	116368.000000	116368.000000
mean	12.190148	23.203007	0.670632	5.093247
std	6.366878	7.085492	1.181365	2.800193
min	-8.500000	-4.800000	0.000000	0.000000
25%	7.700000	18.000000	0.000000	4.000000
50%	12.000000	22.600000	0.000000	4.700000
75%	16.800000	28.200000	0.600000	5.200000
max	31.900000	48.100000	3.200000	21.800000

```
# impute missing values in X_train and X_test with respective column median in X_train
for df1 in [X_train, X_test]:
    for col in numerical:
        col_median=X_train[col].median()
        df1[col].fillna(col_median, inplace=True)
```

```
# check again missing values in numerical variables in X_train
X_train[numerical].isnull().sum()
```

```
mintemp      0
maxtemp      0
rainfall     0
evaporation  0
sunshine     0
windgustspeed 0
windspeed9am 0
windspeed3pm 0
humidity9am   0
humidity3pm   0
pressure9am   0
pressure3pm   0
cloud9am      0
cloud3pm      0
temp9am       0
temp3pm       0
raintoday     0
year          0
month         0
day           0
dtype: int64
```

```
# impute missing categorical variables with most frequent value
for df2 in [X_train, X_test]:
    df2['windgustdir'].fillna(X_train['windgustdir'].mode()[0], inplace=True)
    df2['winddir9am'].fillna(X_train['winddir9am'].mode()[0], inplace=True)
    df2['winddir3pm'].fillna(X_train['winddir3pm'].mode()[0], inplace=True)
```

```
# check missing values in categorical variables in X_train
X_train[categorical].isnull().sum()
```

```
location      0
windgustdir    0
winddir9am     0
winddir3pm     0
dtype: int64
```

```
X_test = pd.concat([X_test[numerical],  
                    pd.get_dummies(X_test.location),  
                    pd.get_dummies(X_test.windgustdir),  
                    pd.get_dummies(X_test.winddir9am),  
                    pd.get_dummies(X_test.winddir3pm)], axis=1)
```



```

results = pd.DataFrame({
    'Model': ['Support Vector Machines', 'KNN', 'Logistic Regression',
              'Random Forest', 'Naive Bayes', 'Perceptron',
              'Stochastic Gradient Decent',
              'Decision Tree'],
    'Score': [acc_linear_svc, acc_knn, acc_log,
              acc_random_forest, acc_gaussian, acc_perceptron,
              acc_sgd, acc_decision_tree]})
result_df = results.sort_values(by='Score', ascending=False)
result_df = result_df.set_index('Score')
result_df.head(9)

```

Model	
Score	
100.00	KNN
100.00	Random Forest
100.00	Decision Tree
84.66	Logistic Regression
84.59	Support Vector Machines
81.05	Stochastic Gradient Decent
73.89	Perceptron
64.37	Naive Bayes

```
# visualize confusion matrix with seaborn heatmap
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'],
                          index=['Predict Positive:1', 'Predict Negative:0'])
sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')

from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred_test))

TP = cm[0,0]
TN = cm[1,1]
FP = cm[0,1]
FN = cm[1,0]
```

	precision	recall	f1-score	support
0.0	0.87	0.95	0.91	22730
1.0	0.73	0.49	0.59	6362
accuracy			0.85	29092
macro avg	0.80	0.72	0.75	29092
weighted avg	0.84	0.85	0.84	29092

